


CS 319 Applied Programming

Fall 2018

Programming Assignment 3

Trees



Instructor: Muhammad Saqib Ilyas

Assigned on: November 15, 2018

Due on: November 28, 2018

Image source: <https://www.bforball.com/drawing-lessons-tree.php>

Summary instructions

Trees are an important data structure used in various applications. In this programming assignment you will first implement (general) trees in C++. Then, you will use the tree data structure to load a dataset of weekly sales by various departments in multiple stores owned by a single retail store chain. You will then use this tree to answer interesting queries on that data.

- Start reading this handout as soon as possible. The more you delay, the deeper you will be in trouble. This assignment will take some time.
- Read this handout carefully before starting implementation.
- Ample starter code has been provided.
- The code in the files `node.h` and `tree.h` define prototypes for the tree node and the tree class itself. The file `helper.h` implements some helper functionality. The code in these three files may not be modified in any way.
- The implementation for the tree node should be done in `node_impl.h`, while the implementation for the tree class should be done in `tree_impl.h`. A skeletal `main()` function is provided to you in the file `pa3.cpp`. You are allowed to make changes to this `main()` function and define any additional functions that you want to call from `main()` in the file `pa3.cpp`, also provided to you.
- Exercise caution with the use of pointer variables. For instance, be careful when calling functions or accessing data members using pointer variables. If the pointer is NULL or invalid, this may cause the program to crash.

Description of individual classes and/or files

Class `TreeNode`

This class defines a node in the tree data structure. It is templated to be generic. It includes pointers to other `TreeNode` objects, specifically, a parent pointer, a pointer to the first child, a pointer to the next sibling and a pointer to the previous sibling. Figure 1 shows the structure of an object of this class. Figure 2 shows how objects to represent a specific tree would be stored. In Figure 2, the value 1 is stored at the root node, whereas the first, second and third child of the root node store values 2, 3, and 4 respectively.




Figure 1: The layout of a `TreeNode` object




Figure 2: An example to show how `TreeNode` objects would store a tree with the node 1 at root and nodes 2, 3 and 4 stored as its children.

An explanation of what each member function of this class is supposed to do, follows. The implementation of these functions must be provided in the file `node_impl.h`:

`TreeNode<T>()`: This is the default constructor. It should assign a default value to all object members.

`TreeNode<T>(T val)`: This is a parameterized constructor. The value in variable `val` should be copied to the value member of the object, while all pointers should be assigned a default value.

`bool IsRoot()`: This function should return true if the current object is the root node of the tree, otherwise return false.

`bool IsExternal()`: This function should return true if the current object is an external node of the tree, otherwise return false.

`int GetHeight()`: This function should return the height of the node in the tree.

`int GetDepth()`: This function should return the depth of the node in the tree.

`T GetValue()`: This function should return the value stored at the current node.

`void SetValue (T val)`: This function should update the value stored at the current node based on what is passed as argument.

`bool AddChild(T val)`: Create a new node object and add it as a child of the current node. For efficiency, the new node should be added as the first child of the current node. If a first child already exists, it should now become the second child. If a child node already has the value in variable `val`, this function should not modify the tree and return false. Return true if the child is successfully added.

`void DeleteChild(T val)`: Delete the child that stores the same value as in the variable `val`, if any.

`TreeNode<T>* SearchChild(T val)`: Search for a child node (NOTE: one level lower only) of this object that stores the value in variable `val`. Return that child node's address. Return NULL if no child node of the current node stores this value.

`TreeNode<T>* GetParent()`: Return the address of the parent node of the current node.

`TreeNode<T>* GetNextSibling()`: Return the address of this node's next sibling.

`TreeNode<T>* GetPrevSibling()`: Return the address of this node's previous sibling.

`TreeNode<T>* GetFirstChild()`: Return the address of the first child of the current node.

Class Tree

This class implements the tree data structure using the class `TreeNode`. This class maintains a pointer to the root node of the tree as the primary means of accessing the nodes in the tree. A secondary means of accessing the tree is the member called `cursor`. This is a pointer, which should initially be the same as `root` and can be moved around over the nodes of the tree using the functions `MoveUp()`, `MoveDown()`, `MoveLeft()`, `MoveRight()` and `ResetCursor()`. Certain member functions of the tree class operate on the node pointed to by `cursor`.

A brief description of what each function in the class Tree is expected to do follows:

`Tree()`: The default constructor. It should create an empty tree. The size member variable should be appropriately assigned.

`~Tree()`: This is the destructor for the tree. This task must be done recursively. For re-use, simply call the `Destroy()` function from the destructor.

`void Deallocate(TreeNode<T>* p)`: Deallocate the subtree rooted at node p. This must be done recursively. Think about whether you would deallocate a tree using pre-order or post-order traversal. Then, apply the appropriate traversal block to free the memory allocated to each node in the subtree rooted at p.

`TreeNode<T>* Root()`: Returns a pointer to the root node of the tree.

`bool IsEmpty()`: Return true if the tree is empty, false otherwise.

`int Size()`: Returns the number of nodes in the tree. This can simply return the value of the size member variable, if you maintain that variable on every add / delete to the tree.

`void PreOrder()`: This is the top level function called to do a pre-order traversal of the tree. This function should call the overloaded `PreOrder()` function with the root node as the first argument and 0 for the second. That function implements the pre-order semantics recursively.

`void PostOrder()`: This is the top level function called to do a post-order traversal of the tree. This function should call the overloaded `PostOrder()` function with the root node as the first argument and 0 as the second. That function implements the post-order semantics recursively.

`void PreOrder(TreeNode<T>* n, int depth)`: This function should do a pre-order traversal of the subtree rooted at the node n. The visit operation is to display the value stored at the node preceded by as many spaces as the value stored in the variable depth. This function must be implemented recursively.

`void PostOrder(TreeNode<T>* n, int depth)`: This function should do a post-order traversal of the subtree rooted at the node n. The visit operation is to display the value stored at the node preceded by as many spaces as the value stored in the variable depth. This function must be implemented recursively.

`TreeNode<T>* Search(T val)`: Searches the entire tree for and returns the pointer to the first node storing the same value as that passed in as argument to this function. If no such node exists, the function should return NULL. The search should be done on the tree nodes in the same order as a pre-order traversal. This function should simply call the function `SearchBelow()` with root node as the first argument.

`TreeNode<T>* SearchBelow(TreeNode<T>* n, T val)`: Searches the subtree rooted at node n for and returns the pointer to the first node storing the same value as that passed in as argument to this function. If no such node exists, the function should return NULL. The search should be done on the tree nodes in the same order as a pre-order traversal.

`TreeNode<T>* SearchChildren(TreeNode<T>*n, T val)`: Searches the immediate children (only one level deep) of the node `n` for a node that stores the value passed in as second argument. If such a child exists, its address is returned. NULL is returned otherwise.

`TreeNode<T>* SearchChildren(T val)`: Searches the immediate children (only one level deep) of the node pointed to by the member variable `cursor`, for a node that stores the value passed in as second argument. If such a child exists, its address is returned. NULL is returned otherwise.

`TreeNode<T>* InsertChild(TreeNode<T>* parent, T val)`: Insert a new child under the node `parent`. The new node should store the value passed in as the second argument to this function. Return the address of the new node. This function should rely on the add child functionality of the `TreeNode` class.

`TreeNode<T>* InsertChild(T val)`: Insert a new child under the node pointed to by the member variable `cursor`. The new node should store the value passed in as the second argument to this function. Return the address of the new node. This function should rely on the add child functionality of the `TreeNode` class.

`void Delete(TreeNode<T>* n, T val)`: Delete the subtree rooted at node `n`. Simply call the `Deallocate()` member function for this.

`void Destroy()`: Deallocate the entire tree. Call the `Deallocate` member function() with the root node as parameter to do this.

`void InsertRootNode(T val)`: If the root node doesn't already exist, insert a new node at the root of the tree, storing value passed in as argument to this function. Otherwise, do nothing.

`bool MoveDown()`: Move the cursor to the first child of the node pointed to by the member variable `cursor` and return true. If the node pointed to by `cursor` has no children, the cursor should not be updated and the function should return false.

`bool MoveRight()`: Move the cursor to the next sibling of the node pointed to by the member variable `cursor` and the function should return true. If the node pointed to by `cursor` has no next sibling, the cursor should not be updated and the function should return false.

`bool MoveUp()`: Move the cursor to the parent of the node pointed to by the member variable `cursor` and return true. If the node pointed to by `cursor` has no parent, the cursor should not be updated and the function should return false.

`bool MoveLeft()`: Move the cursor to the previous sibling of the node pointed to by the member variable `cursor` and the function should return true. If the node pointed to by `cursor` has no previous sibling, the cursor should not be updated and the function should return false.

`TreeNode<T>* GetNode()`: Return the value of the member variable `cursor`.

`void ResetCursor()`: Reset the cursor member variable so that it points to the root node of the tree.

The main() function

The main() function is defined in the file pa3.cpp. It simply calls two test functions, namely test_tiny_tree() and test_tiny_sales().

The test_tiny_tree() function

This function is defined in the file pa3.cpp. In the first three lines, it displays a banner on the screen to show the user what is about to be done. It then creates a tree object to store unsigned int values. It declares three TreeNode pointer variables. The value 1 is then inserted into the empty tree as the root node. The values 2, 3 and 4 are then inserted, respectively into the tree with the root node as the parent, and the pointers to the newly inserted children are stored in pointers p1, p2 and p3, respectively. Then, the values 5 and 6 are inserted as children of the node storing value 2. The values 7 and 8 are inserted as children of the node storing the value 3. The values 9 and 10 are inserted as children of the node storing the value 4. The function then calls the PreOrder() function of the tree class to display the pre-order traversal of the tree, which should look like the output in Figure 3.

```
=====
|           Testing on a tiny tree           |
|=====
1
4
10
9
3
8
7
2
6
5
```




Figure 3: Expected correct output for the test_tiny_tree() function and the graphical representation of the tiny tree

The test_tiny_sales() function

This function first displays a banner on the screen to inform the user what is about to happen. It then creates a tree object. It passes this tree object along with the filename "tiny.csv" to a function ReadFile(). This function is already implemented for you in the file helper.h. The ReadFile() function reads the data file tiny.csv and loads its contents into a tree as described below. The ReadFile() function returns true if the file is successfully loaded into the tree. If not, it returns false. If the file is loaded, test_tiny_sales() function proceeds to call two functions TotalDaySales() and TotalMonthSales() to calculate the total sales for 5/2/2010 and the month of February, 2010, respectively. You must implement these two functions yourself. The test_tiny_sales() function displays the values returned by these functions and terminates.

We have come to possess leaked sales data for a large retail store chain. The file tiny.csv contains a subset of that data. The file's contents are:

Store,Dept,Date,Weekly_Sales,IsHoliday

1,1,2010-02-05,14924.5,FALSE

1,2,2010-02-05,24924.5,FALSE

1,1,2010-02-06,34924.5,FALSE

1,2,2010-02-06,44924.5,FALSE

The first column contains the attribute store ID, the second column stores the attribute department ID, the third column contains the date in YYYY-MM-DD format. The fourth column stores the sales (US \$) made in the week identified by the date specified in the previous column by the store and department identified in the first two columns, respectively. The last column is being ignored.

In the function ReadFile(), we create a root node in the tree with a value 0. This node is pretty much useless, except it pulls all the data together. The children of the root node are all year values encountered in the data file. In this case, there will be only one, storing the value 2010. In a different data file, however, multiple year values may be present, and the root node would, correspondingly, have multiple children.

The function adds month nodes for every month, encountered in the Date column, under the corresponding year node. In the case of tiny.csv, 2010 has only one child with value 2. The function also adds day nodes for every distinct day of the month, encountered in the date column, under the corresponding month node. In the case of tiny.csv, the node 2 has two children, namely 5 and 6. The store IDs become the children of the day nodes and the department IDs become the children of the corresponding store nodes. The actual sales amount for the specific week is inserted as a child of the department node. The sales node is the leaf node. Figure 5 shows a pre-order traversal of the tree loaded as a result of this process using tiny.csv. Also, Figure 6 shows how the data will generally be loaded into a tree level by level.

When this function runs, you should see the following on the screen:

```
=====
|           Testing on sales data           |
=====
File opened!
Calculating total sales for the year 2010 and month 2....
Total sales for year 2010 month 2 and day 5 is 39848
Calculating total sales for the year 2010 and month 2....
Total sales for year 2010 month 2 is 119696
```

Figure 4: Expected correct output from the test_tiny_sales() function

We have not called PreOrder() on the tree here intentionally. The reason is that if you add data to the tiny.csv file to test your code more thoroughly, the pre-order traversal output would be out of control. However, for the purpose of clarity of understanding how the tree is organized by the ReadFile() function, please see Figure 5.

```
0
2010
2
6
1
2
44924
1
34924
5
1
2
24924
1
14924
```

Figure 5: The pre-order traversal of the tiny.csv data




Figure 6: Tree representation of a general dataset as loaded by ReadFile() – partial view. The different levels correspond different attributes of the data.

The TotalDaySales() function

This function should do some basic sanity checking to make sure that nodes for the year, month and day passed in as argument exist in the tree object, also passed in as argument. It should locate the day object node in the tree and pass a pointer to that object to the RecurseTotalDaySales() function. The value returned in that call should be returned by TotalDaySales() function.

The RecurseTotalDaySales() function

This function should calculate and return the total sales in the subtree rooted at the node whose pointer is passed as argument. It must use recursion. Think about whether you would use pre-order or post-order traversal to calculate the total sales under a node. Then, implement that traversal's logic.

The TotalMonthSales() function

This function should do some basic sanity checking to make sure that nodes for the year and month passed in as argument exist in the tree object, also passed in as argument. It should locate the month object node in the tree and pass a pointer to that object to the RecurseTotalMonthSales() function. The value returned in that call should be returned by TotalMonthSales() function.

The RecurseTotalMonthSales() function

This function should calculate and return the total sales in the subtree rooted at the node whose pointer is passed as argument. It must use recursion. Think about whether you would use pre-order or post-order traversal to calculate the total sales under a node. Then, implement that traversal's logic. Should this function be any different from RecurseTotalDaySales() function?

Submission instructions

Once you are done, compress all the .cpp and .h files into one .zip file and upload to slate.nu.edu.pk. Do try to submit before the deadline to avoid any unnecessary hassle. Good luck!